

---

# **Use ROCm™ on Radeon™ GPUs**

## **Documentation**

**Advanced Micro Devices, Inc.**

**Aug 13, 2025**



# CONTENTS

<b>1 Prerequisites to use ROCm on Radeon desktop GPUs for machine learning development</b>	<b>3</b>
1.1 Supported hardware . . . . .	3
1.2 Supported operating systems . . . . .	4
1.3 Recommended system configuration . . . . .	4
1.4 Additional package repositories . . . . .	5
<b>2 How to guides</b>	<b>7</b>
2.1 Linux How to guide - Use ROCm on Radeon GPUs . . . . .	7
2.2 WSL How to guide - Use ROCm on Radeon GPUs . . . . .	32
<b>3 Usecases</b>	<b>43</b>
3.1 vLLM . . . . .	43
3.2 ComfyUI . . . . .	46
<b>4 Compatibility matrices</b>	<b>51</b>
4.1 Linux support matrices by ROCm version . . . . .	51
4.2 WSL support matrices by ROCm version . . . . .	53
<b>5 Limitations and recommended settings</b>	<b>57</b>
5.1 6.4.2 release known issues . . . . .	57
5.2 Multi-GPU configuration . . . . .	58
5.3 Windows Subsystem for Linux (WSL) . . . . .	58
<b>6 AI community</b>	<b>61</b>
<b>7 Report a bug</b>	<b>63</b>



### Turn your desktop into a Machine Learning platform with the latest high-end AMD Radeon™ 9000 and 7000 series GPUs

AMD has expanded support for Machine Learning Development on RDNA™ 4 and RDNA™ 3 GPUs with Radeon™ Software for Linux 25.10.2.1 with ROCm™ 6.4.2.

Researchers and developers working with Machine Learning (ML) models and algorithms using PyTorch, ONNX Runtime, JAX, or TensorFlow can now also use ROCm 6.4.2 on Linux® to tap into the parallel computing power of the newest high-end AMD Radeon GPUs for desktop, including:

- AMD Radeon™ AI PRO R9700
- AMD Radeon™ RX 9070
- AMD Radeon™ RX 9070 XT
- AMD Radeon™ RX 9070 GRE
- AMD Radeon™ RX 9060 XT
- AMD Radeon RX 7700 XT

These new GPUs based on the RDNA 4 architecture join the already-supported Radeon 7000 series built on RDNA 3, further expanding support for high-performance local ML development on Linux®.

A client solution built on powerful high-end AMD GPUs enables a local, private, and cost-effective workflow to develop ROCm and train Machine Learning for users who were solely reliant on cloud-based solutions.

### More ML performance for your desktop

- With today's models easily exceeding the capabilities of standard hardware and software not designed for AI, ML engineers are looking for cost-effective solutions to develop and train their ML-powered applications. Due to the availability of significantly large GPU memory sizes of 24GB or 48GB, utilization of a local PC or workstation equipped with the latest high-end AMD Radeon 9000 & 7000 series GPU offers a robust/potent yet economical option to meet these expanding ML workflow challenges.
- Latest high-end AMD Radeon 9000 series GPUs are built on the RDNA 4 GPU architecture and AMD Radeon 7000 series GPUs are built on the RDNA 3 GPU architecture
  - offers up to 24GB or 48GB of GPU memory to handle large ML models

### Migrate your application from the desktop to the datacenter

- ROCm is the open-source software stack for Graphics Processing Unit (GPU) programming. ROCm spans several domains: General-Purpose computing on GPUs (GPGPU), High Performance Computing (HPC) and heterogeneous computing.
- The latest AMD ROCm 6.4.2 software stack for GPU programming unlocks the massively parallel compute power of these RDNA 4 and RDNA3 GPUs for use with various ML frameworks. The same software stack also supports AMD CDNA™ GPU architecture, so developers can migrate applications from their preferred framework into the datacenter.

### Freedom to customize

ROCM is primarily Open-Source Software (OSS) that allows developers the freedom to customize and tailor their GPU software for their own needs while collaborating with a community of other developers, and helping each other find solutions in an agile, flexible, rapid and secure manner. AMD ROCm allows users to maximize their GPU hardware investment. ROCm is designed to help develop, test and deploy GPU accelerated HPC, AI, scientific computing, CAD, and other applications in a free, open-source, integrated and secure software ecosystem.

### Improved interoperability

- Support for PyTorch, one of the leading ML frameworks.
- Support for ONNX Runtime to perform inference on a wider range of source data, including INT8 with MIGraphX.

- Support for TensorFlow.

#### **Radeon™ Software for Linux® 25.10.2.1 with ROCm 6.4.2 Highlights**

- Support for Llama.cpp
- Forward Attention 2 (FA2) backward pass enablement
- Support for JAX (inference)
- **New models**
  - Llama 3.1
  - Qwen 1.5
  - Qwen 2.5-7B
  - ChatGLM 2/4
  - Deepseek R1 Distilled (7B, 32B, and 70B)
- Support for RedHat Enterprise Linux (RHEL) 9.6 (excluding WSL)

---

**Note:** Visit [AMD ROCm Documentation](#) for the latest on ROCm.

For the latest driver installation packages, visit [Linux Drivers for Radeon Software](#).

---

## PREREQUISITES TO USE ROCM ON RADEON DESKTOP GPUs FOR MACHINE LEARNING DEVELOPMENT

Before starting with the installation, ensure that your system meets the necessary requirements such as supported hardware, a compatible operating system, and the recommended system configuration to ensure optimal performance and functionality.

See [Compatibility matrices](#) for more information.

### 1.1 Supported hardware

#### 1.1.1 Supported graphics processing units

To successfully install ROCm™ for machine learning development, ensure that your system is operating on a Radeon™ Desktop GPU listed in the [Compatibility matrices](#) section.

#### 1.1.2 Recommended memory

The recommended memory to use ROCm on Radeon. These specifications are required for complex AI/ML workloads (for example, large language models):

- 64GB Main Memory
- 24GB GPU Video Memory

**Note** AMD recommends having the same amount of system memory as video memory, as a minimum.

#### Minimum recommendations

Minimum memory requirements to use ROCm on Radeon. Note that low system memory may cause issues running inference models.

- 16GB Main Memory
- 8GB GPU Video Memory

**Important!** These are guidelines only. Note that minimum memory required will vary depending on workload.

## 1.2 Supported operating systems

Ensure that your operating system is up-to-date to successfully install ROCm for machine learning development.

Refer to [Compatibility matrices](#) for up-to-date operating system compatibility.

### 1.2.1 Update Ubuntu® operating system

Use the following commands to bring your OS up-to-date:

```
sudo apt-get update  
sudo apt-get dist-upgrade
```

## 1.3 Recommended system configuration

This section guides users on how to optimize system configuration for ROCm™ usage, ensuring smooth and performant ROCm operation.

### 1.3.1 PCIe atomics for PyTorch

ROCM is an extension of HSA platform architecture, and shares queuing model, memory model, signaling and synchronization protocols.

Platform atomics are integral to perform queuing and signaling memory operations, where there may be multiple-writers across CPU and GPU agents.

For more details, see [How ROCm uses PCIe atomics](#).

### 1.3.2 Disable iGPU

The iGPU is non-essential for AI and ML workloads and not officially supported. Disable iGPU in SBIOS before proceeding to avoid unknown issues.

Alternatively, use environment variables to select the target GPU.

Here are examples to disable iGPU on some AMD motherboards:

#### Gigabyte™ X670 AORUS ELITE AX

1. Enter BIOS

**Path:** *Advanced → AMD CBS → NBIO Common Options → GFX Configuration → iGPU Configuration*

2. Set iGPU to Disabled

## ASUS Prime X670-P WIFI

1. Enter BIOS

**Path:** Advanced → NB Configuration → Integrated Graphics

2. Set to Disabled

**NOTE:** This step only applies to AMD motherboards, no action is required for non-AMD motherboards.

There are no minimum motherboard hardware requirements.

## Alternative option: Use environment variables to select target GPU

An alternative option to disabling the iGPU is to use environment variable to select the GPU.

See [GPU Isolation Techniques](#) to specify the device indices you would like to expose to your application.

## 1.4 Additional package repositories

For some distributions, the ROCm installation packages depend on packages that aren't included in the default package repositories. These external repositories need to be sourced before installation. Use the following instructions specific to your distribution to add the necessary repositories.

### 1.4.1 Red Hat Enterprise Linux

1. Add the EPEL repository for RHEL 9.

```
wget https://dl.fedoraproject.org/pub/epel/epel-release-latest-9.noarch.rpm  
sudo rpm -ivh epel-release-latest-9.noarch.rpm
```

2. Enable the CodeReady Linux Builder (CRB) repository.

In order to enable CRB, you may need to install dnf-plugin-config-manager first.

```
sudo dnf install dnf-plugin-config-manager  
sudo crb enable
```

#### Note

For more information, see [Additional package repositories for Instinct GPUs](#).



## HOW TO GUIDES

These guides walk you through the various installation processes required to pair ROCm™ with the latest high-end AMD Radeon™ 7000 series desktop GPUs.

Linux

WSL

*Linux How to guide*

*WSL How to guide*

### 2.1 Linux How to guide - Use ROCm on Radeon GPUs

This guide walks you through the various installation processes required to pair ROCm™ with the latest high-end AMD Radeon™ 7000 series desktop GPUs, and get started on a fully-functional environment for AI and ML development.

#### 2.1.1 Install Radeon software for Linux with ROCm

The ROCm™ Software Stack and other Radeon™ software for Linux components are installed using the `amdgpu-install` script to assist you in the installation of a coherent set of stack components.

- Simplifies the installation of the AMDGPU stack by encapsulating the distribution specific package installation logic and by using command line options that allows you to specify the:
  - Use case of the AMDGPU stack to be installed (Graphics or Workstation)
  - Combination of components (Pro stack, or user selection)
- Performs post-install checks to verify whether the installation was performed successfully.
- Installs the uninstallation script to allow you to remove the whole AMDGPU stack from the system by using a single command.

The script is provided by the installer package. See [Compatibility matrices](#) for support information.

### Install AMD unified driver package repositories and installer script

Select the applicable Ubuntu® version to download and install the `amdgpu-install` script on the system.

#### Ubuntu 24.04

1. Install kernel headers and development packages.

```
sudo apt install python3-setuptools python3-wheel
```

2. Enter the following commands to install the installer script for Ubuntu version 24.04.

```
sudo apt update
wget https://repo.radeon.com/amdgpu-install/6.4.2.1/ubuntu/noble/amdgpu-install_6.4.
˓→60402-1_all.deb
sudo apt install ./amdgpu-install_6.4.60402-1_all.deb
```

#### Ubuntu 22.04

1. Install kernel headers and development packages.

```
sudo apt install python3-setuptools python3-wheel
```

2. Enter the following commands to install the installer script for Ubuntu version 22.04.

```
sudo apt update
wget https://repo.radeon.com/amdgpu-install/6.4.2.1/ubuntu/jammy/amdgpu-install_6.4.
˓→60402-1_all.deb
sudo apt install ./amdgpu-install_6.4.60402-1_all.deb
```

#### RHEL 9.6

**Note** See [Additional package repositories](#) for more information.

1. Install kernel headers and development packages.

```
sudo dnf install python3-setuptools python3-wheel
```

2. Install the `amdgpu-install` script on the system.

```
sudo yum update
wget https://repo.radeon.com/amdgpu-install/6.4.2.1/rhel/9.6/amdgpu-install-6.4.60402-
˓→1.el9.noarch.rpm
sudo yum install ./amdgpu-install-6.4.60402-1.el9.noarch.rpm
```

## Install AMD unified kernel-mode GPU driver, ROCm, and graphics

After the Unified Driver Deb Package repositories are installed, run the installer script with appropriate --usecase parameters to install the driver components.

AMD recommends installing the Graphics usecase by default. Only consider the alternative install option if you have an applicable Workstation usecase scenario.

Enter the following command to display a list of available usecases:

```
sudo amdgpu-install --list-usecase
```

### Option A: Graphics usecase

AMD recommends installing the Graphics usecase by default.

1. Run the following command to install open source graphics and ROCm.

```
amdgpu-install -y --usecase=graphics,rocm
```

Watch for output warning or errors indicating an unsuccessful driver installation.

**Note:** The `-y` option installs non-interactively. This step may take several minutes, depending on internet connection and system speed.

2. Reboot the system.

```
sudo reboot
```

See [Using the amdgpu-install script](#) for more information.

Next, [set Groups permissions](#).

### Option B: Workstation usecase

Only consider the alternative install option if you have an applicable Workstation usecase scenario.

1. Run the following command to install workstation graphics and ROCm.

```
amdgpu-install -y --usecase=workstation,rocm
```

**Note:** The `-y` option installs non-interactively. This step may take several minutes, depending on internet connection and system speed.

2. Reboot the system.

```
sudo reboot
```

See [Using the amdgpu-install script](#) for more information.

Next, [set Groups permissions](#).

### Set Groups permissions

Once the driver is installed, add any current user to the render and video groups to access GPU resources.

Reboot in order for group changes to take effect.

### Add user to render and video groups

1. Enter the following command to check groups in the system:

```
groups
```

2. Add user to the render and video group using the command:

```
sudo usermod -a -G render,video $LOGNAME
```

3. Reboot the system.

```
sudo reboot
```

See [Setting Permissions for Groups](#) for more information.

### Post-install verification checks

Run these post-installation checks to verify that the installation is complete:

1. Verify that the current user is added to the render and video groups.

```
groups
```

#### Expected result:

```
<username> adm cdrom sudo dip video plugdev render lpadmin lxd sambashare
```

<username> indicates the current user, and this result will vary in your environment.

2. Check if *amdgpu* kernel driver is installed.

```
dkms status
```

#### Expected result:

```
amdgpu/x.x.x-xxxxxxxx.xx.xx, x.x.x-xx-generic, x86_64: installed
```

3. Check if the GPU is listed as an agent.

```
rocminfo
```

#### Expected result:

```
[...]
*****
Agent 2
*****
Name : gfx1100
```

(continues on next page)

(continued from previous page)

<b>Uuid:</b>	GPU-5ecee39292e80c37
Marketing Name:	Radeon RX 7900 XTX
Vendor Name:	AMD
[...]	
[...]	

4. Check if the GPU is listed.

```
clinfo
```

#### Expected result:

[...]	
Platform Name:	AMD Accelerated Parallel Processing
Number of devices:	1
Device Type:	CL_DEVICE_TYPE_GPU
Vendor ID:	1002h
Board name:	Radeon RX 7900 XTX
[...]	

See [Installing the all open usecase](#) for additional troubleshooting tips.

## Advanced install methods

For advanced install methods, such as Multi-Version and Package Manager, refer to [AMD GPU Install Script](#).

## Uninstall ROCm

Run the following command to uninstall the ROCm software stack and other Radeon software for Linux components:

```
sudo amdgpu-uninstall
```

## Upgrade to newer versions of Radeon software for Linux

The recommended method to upgrade is to uninstall, followed by an install.

Radeon Software for Linux does not support in-place upgrades.

### 2.1.2 Install PyTorch for ROCm

Refer to this section for the recommended PyTorch via PIP installation method, as well as Docker-based installation.

#### PCIe atomics

ROCM is an extension of HSA platform architecture, and shares queuing model, memory model, signaling and synchronization protocols.

Platform atomics are integral to perform queuing and signaling memory operations, where there may be multiple-writers across CPU and GPU agents.

For more details, see [How ROCm uses PCIe atomics](#).

### Install methods

AMD recommends the PIP install method to create a PyTorch environment when working with ROCm™ for machine learning development.

Using Docker provides portability, and access to a prebuilt Docker container that has been rigorously tested within AMD. Docker also cuts down compilation time, and should perform as expected without installation issues.

#### Option A: PyTorch via PIP installation

AMD recommends the PIP install method to create a PyTorch environment when working with ROCm™ for machine learning development.

Check [Pytorch.org](#) for latest PIP install instructions and availability. See [Compatibility matrices](#) for support information.

**Note** The latest version of Python module numpy v2.0 is incompatible with the torch wheels for this version. Downgrade to an older version is required. Example: pip3 install numpy==1.26.4

**Note** To install the following wheels, Python 3.12 must be set up.

### Install PyTorch via PIP

1. Enter the following command to unpack and begin set up.

```
sudo apt install python3-pip -y
```

2. Enter this command to update the pip wheel.

```
pip3 install --upgrade pip wheel
```

3. Select the applicable Ubuntu version and enter the commands to install Torch and Torchvision for ROCm AMD GPU support.

This may take several minutes.

**Important!** AMD recommends proceeding with ROCm WHLs available at [repo.radeon.com](#). The ROCm WHLs available at PyTorch.org are not tested extensively by AMD as the WHLs change regularly when the nightly builds are updated.

**Important!** When manually downloading WHLs from [repo.radeon](#), ensure to select the compatible WHLs for specific Python versions.

See [Compatibility matrices](#) for support information.

#### Ubuntu 24.04

```
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torch-2.6.0%2Brocm6.4.2.git76481f7c-cp312-cp312-linux_x86_64.whl
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torchvision-0.21.0%2Brocm6.4.2.git4040d51f-cp312-cp312-linux_x86_64.whl
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/pytorch_triton_rocm-3.2.0%2Brocm6.4.2.git7e948ebf-cp312-cp312-linux_x86_64.whl
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torchaudio-2.6.0%2Brocm6.4.2.gitd8831425-cp312-cp312-linux_x86_64.whl
```

(continues on next page)

(continued from previous page)

```
pip3 uninstall torch torchvision pytorch-triton-rocm
pip3 install torch-2.6.0+rocm6.4.2.git76481f7c-cp312-cp312-linux_x86_64.whl_
↳torchvision-0.21.0+rocm6.4.2.git4040d51f-cp312-cp312-linux_x86_64.whl_
↳torchaudio-2.6.0+rocm6.4.2.gitd8831425-cp312-cp312-linux_x86_64.whl pytorch_
↳triton_rocm-3.2.0+rocm6.4.2.git7e948ebf-cp312-cp312-linux_x86_64.whl
```

**Note**

The `--break-system-packages` flag must be added when installing wheels for Python 3.12 in a non-virtual environment.

**Ubuntu 22.04**

```
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torch-2.6.0%2Brocm6.4.
↳2.git76481f7c-cp310-cp310-linux_x86_64.whl
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torchvision-0.21.0
↳%2Brocm6.4.2.git4040d51f-cp310-cp310-linux_x86_64.whl
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/pytorch_triton_rocm-3.
↳2.0%2Brocm6.4.2.git7e948ebf-cp310-cp310-linux_x86_64.whl
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torchaudio-2.6.0
↳%2Brocm6.4.2.gitd8831425-cp310-cp310-linux_x86_64.whl
pip3 uninstall torch torchvision pytorch-triton-rocm
pip3 install torch-2.6.0+rocm6.4.2.git76481f7c-cp310-cp310-linux_x86_64.whl_
↳torchvision-0.21.0+rocm6.4.2.git4040d51f-cp310-cp310-linux_x86_64.whl_
↳torchaudio-2.6.0+rocm6.4.2.gitd8831425-cp310-cp310-linux_x86_64.whl pytorch_
↳triton_rocm-3.2.0+rocm6.4.2.git7e948ebf-cp310-cp310-linux_x86_64.whl
```

**RHEL 9.6**

```
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torch-2.6.0%2Brocm6.4.
↳2.git76481f7c-cp312-cp312-linux_x86_64.whl
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torchvision-0.21.0
↳%2Brocm6.4.2.git4040d51f-cp312-cp312-linux_x86_64.whl
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/pytorch_triton_rocm-3.
↳2.0%2Brocm6.4.2.git7e948ebf-cp312-cp312-linux_x86_64.whl
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torchaudio-2.6.0
↳%2Brocm6.4.2.gitd8831425-cp312-cp312-linux_x86_64.whl
pip3 uninstall torch torchvision pytorch-triton-rocm
pip3 install torch-2.6.0+rocm6.4.2.git76481f7c-cp312-cp312-linux_x86_64.whl_
↳torchvision-0.21.0+rocm6.4.2.git4040d51f-cp312-cp312-linux_x86_64.whl_
↳torchaudio-2.6.0+rocm6.4.2.gitd8831425-cp312-cp312-linux_x86_64.whl pytorch_
↳triton_rocm-3.2.0+rocm6.4.2.git7e948ebf-cp312-cp312-linux_x86_64.whl
```

**Note**

The `--break-system-packages` flag must be added when installing wheels for Python 3.12 in a non-virtual environment.

Next, *verify your PyTorch installation.*

### Option B: Docker installation

Using Docker provides portability, and access to a prebuilt Docker container that has been rigorously tested within AMD. Docker also cuts down compilation time, and should perform as expected without installation issues.

#### Prerequisites to install PyTorch using Docker

Docker for Ubuntu® must be installed.

To install Docker for Ubuntu, enter the following command:

```
sudo apt install docker.io
```

#### Use Docker image with pre-installed PyTorch

Follow these steps to install using a Docker image.

**Note** The latest version of Python module numpy v2.0 is incompatible with the torch wheels for this version.

Downgrade to an older version is required. Example: pip3 install numpy==1.26.4

1. Select the applicable Ubuntu version and enter the following command to pull the public PyTorch Docker image.

**Optional:** You can also download a specific and supported configuration with different user-space ROCm versions, PyTorch versions, and supported operating systems.

Refer to [hub.docker.com/r/rocm/pytorch](https://hub.docker.com/r/rocm/pytorch) to download the PyTorch Docker image.

#### Ubuntu 24.04

```
sudo docker pull rocm/pytorch:rocm6.4.2_ubuntu24.04_py3.12_pytorch_release_2.6.0
```

#### Ubuntu 22.04

```
sudo docker pull rocm/pytorch:rocm6.4.2_ubuntu22.04_py3.10_pytorch_release_2.6.0
```

2. Select the applicable Ubuntu version and start a Docker container using the downloaded image.

#### Ubuntu 24.04

```
sudo docker run -it \
--cap-add=SYS_PTRACE \
--security-opt seccomp=unconfined \
--device=/dev/kfd \
--device=/dev/dri \
--group-add video \
--ipc=host \
--shm-size 8G \
rocm/pytorch:rocm6.4.2_ubuntu24.04_py3.12_pytorch_release_2.6.0
```

## Ubuntu 22.04

```
sudo docker run -it \
--cap-add=SYS_PTRACE \
--security-opt seccomp=unconfined \
--device=/dev/kfd \
--device=/dev/dri \
--group-add video \
--ipc=host \
--shm-size 8G \
rocm/pytorch:rocm6.4.2_ubuntu22.04_py3.10_pytorch_release_2.6.0
```

This will automatically download the image if it does not exist on the host. You can also pass the `-v` argument to mount any data directories from the host onto the container.

Next, [verify the PyTorch installation](#).

See PyTorch Installation for ROCm for more information.

## Verify PyTorch installation

Confirm if PyTorch is correctly installed.

1. Verify if Pytorch is installed and detecting the GPU compute device.

```
python3 -c "import torch" 2> /dev/null && echo "Success" || echo "Failure"
```

**Expected result:**

Success

2. Enter command to test if the GPU is available.

```
python3 -c "import torch; print(torch.cuda.is_available())"
```

**Expected result:**

True

3. Enter command to display installed GPU device name.

```
python3 -c "import torch; print(f'device name [0]:', torch.cuda.get_device_
name(0))"
```

**Expected result:** Example: *device name [0]: Radeon RX 7900 XTX*

```
device name [0]: <Supported AMD GPU>
```

4. Enter command to display component information within the current PyTorch environment.

```
python3 -m torch.utils.collect_env
```

**Expected result:**

```
PyTorch version
ROCM used to build PyTorch
```

(continues on next page)

(continued from previous page)

OS  
Is CUDA available  
GPU model and configuration  
HIP runtime version  
MIOpen runtime version

Environment set-up is complete, and the system is ready for use with PyTorch to work with machine learning models, and algorithms.

### 2.1.3 Install ONNX Runtime for Radeon GPUs

#### Overview

Ensure that the following prerequisite installations are successful before proceeding to install ONNX Runtime for use with ROCm™ on Radeon™ GPUs.

#### Prerequisites

- Radeon Software for Linux (with ROCm) is installed.
- *MIGraphX* is installed. This enables ONNX Runtime to build the correct MIGraphX Execution Provider (EP).
- The half library is installed. See *Verify if MIGraphX is installed with the half library*.
- If the prerequisite installations are successful, proceed to *install ONNX Runtime*.

**NOTE** Unless adding custom features, use the pre-built Python wheel files provided in the PIP installation method.

#### Verify MIGraphX installation

1. Verify if MIGraphX is installed with the half library

```
$ dpkg -l | grep migraphx
$ dpkg -l | grep half
```

#### Expected result:

```
root@aus-navi3x-02:/workspace/AMDMIGraphX# dpkg -l | grep migraphx
ii  migraphx                  2.9.0          amd64
  ↳AMD's graph optimizer
ii  migraphx-dev               2.9.0          amd64
  ↳AMD's graph optimizer
ii  migraphx-tests             2.9.0          amd64
  ↳AMD's graph opt

$ dpkg -l | grep half
```

(continues on next page)

(continued from previous page)

ii half	1.12.0.60000~91~20.04	amd64	
↳ HALF-PRECISION FLOATING POINT LIBRARY			

**Note** Versions may vary between ROCm builds and installed versions of MIGraphX, but the desired result is the same.

2. The half library should come packaged with MIGraphX. If not, it can be installed with the following command.

```
sudo apt install half
```

3. Perform a simple inference with MIGraphX to verify the installation.

```
/opt/rocm-6.4.2/bin/migraphx-driver perf --test
```

## Install ONNX Runtime

### Important!

- Use the provided pre-built Python wheel files from the PIP installation method, unless adding custom features.
- The wheel file contains the MIGraphX and ROCm Execution Providers (EP). Refer to [Install MIGraphX for ONNX RT](#) for more information.
- Refer to [ONNX Runtime Documentation](#) for additional information on ONNX Runtime topics.
- See [ONNX Runtime Tutorials](#) to try out real applications and tutorials on how to get started.

### Option A: PIP install (Recommended)

#### Option A: ONNX Runtime install via PIP installation method (Recommended)

AMD recommends the PIP install method to create an ONNX Runtime environment when working with ROCm for machine learning development.

**Note** The latest version of Python module numpy v2.0 is incompatible with the ONNX Runtime wheels for this version. Downgrade to an older version is required. Example: pip3 install numpy==1.26.4

#### To install via PIP,

Enter this command to download and install the ONNX Runtime wheel.

```
pip3 uninstall onnxruntime-rocm
pip3 install onnxruntime-rocm -f https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.
↳ 2/
```

Next, *verify your ONNX Runtime installation*.

### Option B: Build from source (Advanced)

#### Option B: Build from source for your environment, followed by local wheel file installation (Advanced)

Use this method for advanced customization usecases. This requires the user install the desired ROCm and MIGraphX versions, and creation of softlink prior to starting the build.

**NOTE** The build time typically takes ~45 minutes.

#### Prerequisites to build ONNX from source

- Radeon Software for Linux (with ROCm) is installed
- *MigraphX* is installed
- Softlink is created

To create a softlink for /opt/rocm, enter the following command:

```
language = bash
linenumbers = true
ln -s /opt/rocm* /opt/rocm
```

#### To build from source,

1. Clone the ONNX Runtime repository into the root directory.

```
cd /
git clone https://github.com/microsoft/ONNXRuntime.git
```

2. Git clone AMDMIGraphX into the home folder.

```
cd ~
git clone https://github.com/ROCm/AMDMIGraphX.git
```

3. Create a docker image for MIGraphX.

#### Note

- Refer to [AMDMIGraphX Github](#) for up-to-date ONNX Runtime and MIGraphX dependencies.
- MIGraphX can still be built or installed from apt.
- For MIGraphX package builds via **RBuild**, refer to [these MIGraphX Github instructions](#) to build within the docker container environment.
- For MIGraphX package builds via **CMake**, refer to [these MIGraphX Github instructions](#) to build within the docker container environment.
- Use the `groups` command to ensure that the user is part of the video, render, and docker groups in Linux to run the docker container.

```
groups
tthemist@aus-navi3x-02 ~/
tthemist sudo video render docker
```

- Run the following for a simple MIGraphX apt install:

```
cd AMDMIGraphX
docker build -t migraphx .
docker run --device=/dev/kfd --device=/dev/dri -v=/tmp:/code/AMDMIGraphX -v=/onnxruntime:/onnxruntime -w /code/AMDMIGraphX --group-add video -it --migraphx
apt install migraphx migraphx-dev half
```

- Run rocm-smi to ensure that ROCm is installed and detects the supported GPU(s).

```
$ rocm-smi
```

**Expected result:**

```
=====
ROCM System Management Interface
=====
=====
Concise Info
=====
Device [Model : Revision] Temp Power Partitions SCLK MCLK Fan
Perf PwrCap VRAM% GPU%
Name (20 chars) (Edge) (Avg) (Mem, Compute)
=====
0 [0x0e0d : 0x00] 32.0°C 73.0W N/A, N/A 1526Mhz 96Mhz 31.76
auto 241.0W 0% 50%
0x7448
=====
End of ROCm SMI Log
=====
```

- Configure Git to treat all directories as safe to use and run the build script.

```
cd AMDMIGraphX
git config --global --add safe.directory **
/tools/build_and_test_onnxrt.sh
```

This builds ONNX Runtime and adds ROCm and MIGraphX EP support to the ONNX Runtime interface and requires multiple external repo pieces be checked out automatically prior to the build.

- Install ONNX Runtime once MIGraphX is built.

```
$ pip3 install /onnxruntime/build/Linux/Release/dist/*.whl
```

Next, *verify your ONNX Runtime installation.*

## Verify ONNX Runtime installation

Verify that the install works correctly by performing a simple inference with MIGraphX.

```
python3 -c "import onnxruntime as ort; print(ort.get_available_providers())"
```

**Expected result:** The following EPs are displayed.

```
>>> import onnxruntime as ort
>>> ort.get_available_providers()
['MIGraphXExecutionProvider', 'ROCMExecutionProvider', 'CPUExecutionProvider']
```

This indicates that the `MIGraphXExecutionProvider` and `ROCMExecutionProvider` are now running on the system, and the proper ONNX Runtime package has been installed.

Installation is complete and ONNX Runtime is available through the Python interface library, as well as scripts that invoke ONNX Runtime inference sessions.

For more information on the ONNX Runtime Python library, refer to [Get started with ONNX Runtime in Python](#).

### 2.1.4 Install TensorFlow for ROCm

TensorFlow is an open-source library for solving machine-learning, deep-learning, and artificial-intelligence problems. It can be used to solve many problems across different sectors and industries but primarily focuses on training and inference in neural networks. It is one of the most popular and in-demand frameworks and is very active in open source contribution and development.

**Important!** ROCm on Radeon scripts uses Keras 2, but newest wheels uses Keras 3 by default.

A manual install of the `tf-keras` package is required to enable Keras 2 on TensorFlow.

Use the following command to install the ROCm compatible TensorFlow wheel.

```
pip install tf-keras --no-deps
```

**Additional information** As of ROCm 6.1, tensorflow-rocm packages are found at <https://repo.radeon.com/rocm/manylinux>. Prior to ROCm 6.1, packages were found at <https://pypi.org/project/tensorflow-rocm>. Refer to the following version support matrix:

ROCM version	TensorFlow version
6.1.x	2.13.1, 2.14.0, 2.15.0
6.0.x	2.12, 2.13.1, 2.14.0

### Pre-requisites

- Radeon software for Linux (with ROCm) must be installed.
- MIGraphX must be installed for TensorFlow to build the correct mig execution provider.

### PIP installation

Use the PIP install method to create a TensorFlow environment when working with ROCm for machine learning development.

**Note** The latest version of Python module numpy v2.0 is incompatible with the TensorFlow wheels for this version. Downgrade to an older version is required. Example: `pip3 install numpy==1.26.4`

**Library Compatibility** When installing TensorFlow, it is essential to ensure that any additional TensorFlow-related libraries, or dependencies are compatible with the version of TensorFlow that you are using.

TensorFlow libraries (such as tensorflow-hub, etc.) often have specific version requirements that depend on the main TensorFlow version.

Be aware that installing or upgrading TensorFlow libraries may inadvertently replace the ROCm-supported TensorFlow package with a non-ROCM supported version.

To avoid this, ensure that you are explicitly specifying the ROCm-compatible version during installation.

To install TensorFlow,

### Ubuntu 24.04

Download and install the TensorFlow wheel.

#### Note

The `--break-system-packages` flag must be added when installing wheels for Python 3.12 in a non-virtual environment.

```
pip3 uninstall tensorflow-rocm
pip3 install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/tensorflow_rocm-2.
˓→18.1-cp312-cp312-manylinux_2_28_x86_64.whl
```

### Ubuntu 22.04

Download and install the TensorFlow wheel.

```
pip3 uninstall tensorflow-rocm
pip3 install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/tensorflow_rocm-2.
˓→18.1-cp310-cp310-manylinux_2_28_x86_64.whl
```

### RHEL 9.6

#### Note

The `--break-system-packages` flag must be added when installing wheels for Python 3.12 in a non-virtual environment.

```
pip3 uninstall tensorflow-rocm
pip3 install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/tensorflow_rocm-2.
˓→18.1-cp312-cp312-manylinux_2_28_x86_64.whl
```

Next, *verify your TensorFlow installation.*

### Verify TensorFlow installation

To test the TensorFlow installation, run the container image as specified in the previous section [Installing TensorFlow](#). Ensure you have access to the Python shell in the Docker container.

```
python3 -c 'import tensorflow' 2> /dev/null && echo 'Success' || echo 'Failure'
```

Next, *run basic TensorFlow example.*

## Run basic TensorFlow example

The TensorFlow examples repository provides basic examples that exercise the framework's functionality.

The MNIST database is a collection of handwritten digits that may be used to train a Convolutional Neural Network for handwriting recognition.

This dataset is included with your TensorFlow installation.

1. Run the following sample code to load the MNIST dataset, then train and evaluate it.

```
import tensorflow as tf
print("TensorFlow version:", tf.__version__)
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10)
])
predictions = model(x_train[:1]).numpy()
tf.nn.softmax(predictions).numpy()
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
loss_fn(y_train[:1], predictions).numpy()
model.compile(optimizer='adam',
              loss=loss_fn,
              metrics=['accuracy'])
model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test, verbose=2)
```

2. If successful, you should see the following output indicating the image classifier is now trained to around 98% accuracy on this dataset.

```
I0000 00:00:1716848656.190857    212 device_compiler.h:186] Compiled cluster using XLA! This line is logged at
fetime of the process.
1858/1875 [=====>.] - ETA: 0s - loss: 0.2965 - accuracy: 0.91442024-05-27 22:24:18.503331
common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
1875/1875 [=====] - 3s 924us/step - loss: 0.2952 - accuracy: 0.9148
2024-05-27 22:24:18.504658: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
Epoch 2/5
1875/1875 [=====] - 2s 923us/step - loss: 0.1439 - accuracy: 0.9571
Epoch 3/5
1875/1875 [=====] - 2s 924us/step - loss: 0.1098 - accuracy: 0.9668
Epoch 4/5
1875/1875 [=====] - 2s 924us/step - loss: 0.0900 - accuracy: 0.9724
Epoch 5/5
1875/1875 [=====] - 2s 925us/step - loss: 0.0762 - accuracy: 0.9760
2024-05-27 22:24:25.470030: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.470639: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.473927: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.474474: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.479210: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.479879: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.487077: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.487564: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.490836: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.491442: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.491931: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.492460: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.494832: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.497105: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.499508: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.501408: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
2024-05-27 22:24:25.613366: I tensorflow/core/common_runtime/gpu_fusion_pass.cc:508] ROCm Fusion is enabled.
```

Environment set-up is complete, and the system is ready for use with TensorFlow to work with machine learning models, and algorithms.

## 2.1.5 Install Triton for ROCm

OpenAI has developed a powerful GPU focused programming language and compiler called Triton that works seamlessly with AMD GPUs. The goal of Triton is to enable AI engineers and scientists to write high-performant GPU code with minimal expertise.

Triton kernels are performant because of their blocked program representation, allowing them to be compiled into highly optimized binary code. Triton also leverages Python for kernel development, making it both familiar and accessible.

The kernels can be compiled by declaring the `triton.jit` python decorator before the kernel.

### Pre-requisites

- Compatible AMD GPU
  - Linux and ROCm 5.7+ is installed
- See [Compatibility matrices](#) for support information.

### Install libraries

If ROCm 6.0 and the latest version of PyTorch is not installed, the required libraries must first be installed. However, if you encounter issues running any of the commands, we recommend updating with the nightly wheels. This will also install the version of Triton that is compatible with PyTorch for ROCm.

1. Enter the following command to install the libraries.

```
pip install matplotlib pandas -q
pip install --pre torch torchvision torchaudio --index-url https://download.
→pytorch.org/whl/nightly/rocm6.0/ -q
```

2. Enter the following command to import the libraries.

```
import torch
import triton
import triton.language as tl
```

Now, a Triton kernel that approximates the GELU (Gaussian Error Linear Unit) kernel using tanh can be developed.

For more information on how to develop a kernel for GELU and benchmark its performance with its PyTorch analogues, see [Developing Triton Kernels on AMD GPUs](#).

## 2.1.6 Install JAX for ROCm

This directory provides setup instructions and necessary files to build, test, and run JAX with ROCm support using a PIP or Docker install, suitable for both runtime and CI workflows.

**Note** These instructions are for JAX installation on **Radeon** GPUs. To install ROCm on **Instinct** GPUs, refer to [ROCM Instinct documentation](#).

### Install JAX

Follow these instructions to install JAX via PIP or Docker install.

#### PIP installation

Follow these instructions to install JAX via PIP.

**Important** The packages must be installed in the following order:

1. Install pjrt wheel.
2. Install plugin wheel.
3. Install jaxlib wheel.
4. Install jax wheel.

#### Ubuntu 24.04

Install JAX for Ubuntu 24.04.

1. Uninstall previous version

```
pip3 uninstall -y jax-rocm60-pjrt jax-rocm60-plugin jaxlib jax
```

2. Install the pjrt wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax_rocm60_pjrt-  
→0.4.35-py3-none-manylinux_2_28_x86_64.whl
```

3. Install the plugin wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax_rocm60_-  
→plugin-0.4.35-cp312-cp312-manylinux_2_28_x86_64.whl
```

4. Install the jaxlib wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jaxlib-0.4.35-  
→cp312-cp312-manylinux_2_28_x86_64.whl
```

5. Install the jax wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax-0.4.35-py3-  
→none-any.whl
```

## Ubuntu 22.04

Install JAX for Ubuntu 22.04.

1. Uninstall previous version

```
pip3 uninstall -y jax-rocm60-pjrt jax-rocm60-plugin jaxlib jax
```

2. Install the pjrt wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax_rocm60_pjrt-
→0.4.35-py3-none-manylinux_2_28_x86_64.whl
```

3. Install the plugin wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax_rocm60_
→plugin-0.4.35-cp310-cp310-manylinux_2_28_x86_64.whl
```

4. Install the jaxlib wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jaxlib-0.4.35-
→cp310-cp310-manylinux_2_28_x86_64.whl
```

5. Install the jax wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax-0.4.35-py3-
→none-any.whl
```

## RHEL 9.6

Install JAX for RHEL 9.6.

1. Uninstall previous version

```
pip3 uninstall -y jax-rocm60-pjrt jax-rocm60-plugin jaxlib jax
```

2. Install the pjrt wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax_rocm60_pjrt-
→0.4.35-py3-none-manylinux_2_28_x86_64.whl
```

3. Install the plugin wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax_rocm60_
→plugin-0.4.35-cp312-cp312-manylinux_2_28_x86_64.whl
```

4. Install the jaxlib wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jaxlib-0.4.35-
→cp312-cp312-manylinux_2_28_x86_64.whl
```

5. Install the jax wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax-0.4.35-py3-
→none-any.whl
```

### Docker installation

The ROCm JAX team provides prebuilt Docker images, which is the simplest way to use JAX on ROCm. These images are available on Docker Hub and come with JAX configured for ROCm.

#### Ubuntu 24.04

Install JAX for Ubuntu 24.04.

1. To pull the latest ROCm JAX Docker image, run:

```
docker pull rocm/jax:rocm6.4.2-jax0.4.35-py3.12
```

**Note** For specific versions of JAX, review the periodically pushed Docker images at ROCm JAX Community on Docker Hub.

Additional Docker images are available at ROCm JAX on Docker Hub. These contain the latest ROCm version but might use an older version of JAX.

2. Once the image is downloaded, launch a container using the following command:

```
docker run -it -d --network=host --device=/dev/kfd --device=/dev/dri --ipc=host --  
--shm-size 64G \  
--group-add video --cap-add=SYS_PTRACE --security-opt seccomp=unconfined -v  
--$pwd :/jax_dir \  
--name rocm_jax rocm/jax-community:latest /bin/bash  
  
docker attach rocm_jax
```

**Note** The `--shm-size` parameter allocates shared memory for the container. Adjust it based on your system's resources if needed. Replace `$(pwd)` with the absolute path to the directory you want to mount inside the container. If you prefer to use `rocm/jax`, remember to replace `rocm/jax-community` with `rocm/jax`.

#### Ubuntu 22.04

Install JAX for Ubuntu 22.04.

1. To pull the latest ROCm JAX Docker image, run:

```
docker pull rocm/jax:rocm6.4.2-jax0.4.35-py3.10
```

**Note** For specific versions of JAX, review the periodically pushed Docker images at ROCm JAX Community on Docker Hub.

Additional Docker images are available at ROCm JAX on Docker Hub. These contain the latest ROCm version but might use an older version of JAX.

2. Once the image is downloaded, launch a container using the following command:

```
docker run -it -d --network=host --device=/dev/kfd --device=/dev/dri --ipc=host --  
--shm-size 64G \  
--group-add video --cap-add=SYS_PTRACE --security-opt seccomp=unconfined -v  
--$pwd :/jax_dir \  
--name rocm_jax rocm/jax-community:latest /bin/bash  
  
docker attach rocm_jax
```

**Note** The `--shm-size` parameter allocates shared memory for the container. Adjust it based on your system's resources if needed. Replace `$(pwd)` with the absolute path to the directory you want to mount inside the container. If you prefer to use rocm/jax, remember to replace rocm/jax-community with rocm/jax.

## RHEL 9.6

Install JAX for RHEL 9.6.

1. To pull the latest ROCm JAX Docker image, run:

```
docker pull rocm/jax:rocm6.4.2-jax0.4.35-py3.12
```

**Note** For specific versions of JAX, review the periodically pushed Docker images at ROCm JAX Community on Docker Hub.

Additional Docker images are available at ROCm JAX on Docker Hub. These contain the latest ROCm version but might use an older version of JAX.

2. Once the image is downloaded, launch a container using the following command:

```
docker run -it -d --network=host --device=/dev/kfd --device=/dev/dri --ipc=host --  
--shm-size 64G \  
--group-add video --cap-add=SYS_PTRACE --security-opt seccomp=unconfined -v  
$(pwd):/jax_dir \  
--name rocm_jax rocm/jax-community:latest /bin/bash  
  
docker attach rocm_jax
```

**Note** The `--shm-size` parameter allocates shared memory for the container. Adjust it based on your system's resources if needed. Replace `$(pwd)` with the absolute path to the directory you want to mount inside the container. If you prefer to use rocm/jax, remember to replace rocm/jax-community with rocm/jax.

## Verify installation

Refer to [Testing your JAX installation with ROCm](#) for verification steps.

### 2.1.7 Install MIGraphX for Radeon GPUs

MIGraphX is AMD's graph inference engine that accelerates machine learning model inference, and can be used to accelerate workloads within the Torch MIGraphX and ONNX Runtime backend frameworks.

- Torch-MIGraphX, which integrates MIGraphX with PyTorch
- MIGraphX for ONNX Runtime backend, which integrates MIGraphX with ONNX

ONNX Runtime can be driven by either the ROCm™ Execution Provider (EP) or MIGraphX EP

### Introduction to MIgraphX

MIgraphX is a graph optimizer that accelerates the inference for deep learning models. It provides C++ and Python APIs that are integrated within frameworks like Torch MIgraphX, ONNX Runtime, and other user solutions. The following process summarizes the procedures that occur under-the-hood during the optimization and real-time compilation process.

MIgraphX accelerates the Machine Learning models by leveraging several graph-level transformations and optimizations. These optimizations include:

- Operator fusion
- Arithmetic simplifications
- Dead-code elimination
- Common subexpression elimination (CSE)
- Constant propagation

When the aforementioned optimizations are applied, MIgraphX emits code for the AMD GPU by calling to MIOpen, rocBLAS, or creating HIP kernels for a particular operator. MIgraphX can also target CPUs using DNNL or ZenDNN libraries.

For more information on how to install MIgraphX, refer to [AMD MIgraphX Github](#).

### Prerequisites

- Radeon™ Software for Linux (with ROCm) is installed

### Install MIgraphX

Install MIgraphX on your computer. Once the install is completed and verified, proceed to [install Torch-MIgraphX](#) or [MIgraphX for ONNX Runtime](#).

Run the following command to install MIgraphX:

```
$ sudo apt install migraphx
```

Next, proceed to install [Torch-MIgraphX](#) or [MIgraphX for ONNX Runtime](#) as applicable.

### Install Torch-MIgraphX

Install Torch-MIgraphX using the Docker installation method, or build from source.

#### Option A: Docker installation

Using Docker provides portability, and access to a prebuilt Docker container that has been rigorously tested within AMD. Docker also cuts down compilation time, and should perform as expected without installation issues.

1. Clone the `torch_migraphx` repository.

```
git clone https://github.com/ROCMSoftwarePlatform/torch_migraphx.git
```

2. Change directory to `torch-migraphx`.

```
cd torch_migraphx/
```

3. Build image using the provided script.

```
sudo ./build_image.sh
```

4. Run container.

```
sudo docker run -it --network=host --device=/dev/kfd --device=/dev/dri --group-add=video --ipc=host --cap-add=SYS_PTRACE --security-opt seccomp=unconfined --torch_migraphx
```

Next, *verify the Torch-MIGraphX installation*.

### Option B: Build from source

To build from source in a custom environment, refer to the [torch\\_migraphx](#) repository for build steps.

Next, *verify the Torch-MIGraphX installation*.

### Verify Torch-MIGraphX installation

Verify if the Torch-MIGraphX installation is successful.

1. Verify if torch\_migraphx can be imported as a Python module.

```
python3 -c 'import torch_migraphx' 2> /dev/null && echo 'Success' || echo 'Failure'
```

2. Run unit tests.

```
pytest ./torch_migraphx/tests
```

Installation is complete and the system is able to run PyTorch through the python interface library, and scripts that invoke PyTorch inference sessions.

### Install and verify MIGraphX for ONNX Runtime

See [Install ONNX Runtime for Radeon GPUs](#) for MIGraphX for ONNX Runtime installation and verification instructions.

## 2.1.8 mGPU setup and configuration

### Hardware and software considerations

Refer to the following hardware and software considerations to ensure optimal performance.

### Hardware considerations

- **PCIe® slots** AMD recommends a system with multiple x16 (Gen 4) slots, with optimal performance achieved by provision of a 1:1 ratio between the number of x16 slots and the number of GPUs used.  
**Note** Functionality is maintained in the instance where only one x16 slot is available, at the cost of some performance.
- **mGPU power setup** MultiGPU configurations require adequate amounts of power for all the components required. Consult [AMD Radeon™ RX](#) or [AMD Radeon™ PRO](#) for GPU specifications and graphics card power requirements.

### Software considerations

There are no differences in software requirements between single-GPU and multi-GPU usage.

#### mGPU configuration by framework

**Note** PyTorch, ONNX, and Tensorflow may have additional guidelines regarding mGPU configuration. Refer to official mGPU support documentation of the applicable framework for more information.

##### PyTorch

###### Warning

It is recommended to use `DistributedDataParallel`, instead of `DataParallel` to do multi-GPU training, even if there is only a single node. See: [Use nn.parallel.DistributedDataParallel instead of multiprocessing or nn.DataParallel and Distributed Data Parallel](#).

#### mGPU known issues and limitations

AMD has identified common errors when running ROCm™ on Radeon™ multi-GPU configuration at this time, along with the applicable recommendations.

#### IOMMU limitations and guidance

For any issues with application hangs, or problems running a workload when running on a system with multiple GPUs, see [Issue #5: Application hangs on Multi-GPU systems](#).

#### Windows Subsystem for Linux (WSL) support

Microsoft does not currently support mGPU setup in WSL.

## Simultaneous parallel compute workloads

Radeon GPUs does not support large amounts of simultaneous, parallel workloads. It is not recommended to exceed 2 simultaneous compute workloads, with the assumption that workloads are running alongside a graphics environment (eg: Linux desktop).

## Recommended multi-GPU system configuration

PCIe slots connected to the GPU must have identical PCIe lane width or bifurcation settings, and support PCIe 3.0 Atomics.

Refer to [How ROCm uses PCIe Atomics](#) for more information.

### Example:

- ✓ - GPU0 PCIe x16 connection + GPU1 PCIe x16 connection
- ✓ - GPU0 PCIe x8 connection + GPU1 PCIe x8 connection
- X - GPU0 PCIe x16 connection + GPU1 PCIe x8 connection

### Important!

- Only use PCIe slots connected by the CPU and to avoid PCIe slots connected via chipset. Refer to product-specific motherboard documentation for PCIe electrical configuration.
- Ensure the system Power Supply Unit (PSU) has sufficient wattage to support multiple GPUs.

## GPU isolation techniques

For more information, see [GPU isolation techniques](#).

## PCIe atomic operations

Some consumer grade motherboards may only support the first PCIe slot. For unexpected issues, see [How ROCm uses PCIe atomics](#).

## Errors due to GPU and PCIe configuration

When using two AMD Radeon 7900XTX GPUs, the following HIP error is observed when running PyTorch micro-benchmarking if any one of the two GPUs are connected to a non-CPU PCIe slot (PCIe on chipset):

```
RuntimeError: HIP error: the operation cannot be performed in the present state
HIP kernel errors might be asynchronously reported at some other API call, so the
→stacktrace below might be incorrect.
For debugging consider passing HIP_LAUNCH_BLOCKING=1.
Compile with TORCH_USE_HIP_DSA to enable device-side assertions.
```

## Potential GPU reset with some mixed graphics and compute workloads

Working with certain mixed graphics and compute workloads may result in a GPU reset on Radeon GPUs.

Currently identified scenarios include:

- Running multiple ML workloads simultaneously while using the desktop
- Running ML workloads while simultaneously using Blender/HIP

## Accelerate Library & Multiple GPU devices

### Avoid unexpected out of memory errors when working with large models & inference

Refer to the following sections on `infer_auto_device_map()` to explicitly define max memory for each GPU device, when working with accelerate library & multiple GPU devices.

- Big model inference
- Limits and further development

## 2.2 WSL How to guide - Use ROCm on Radeon GPUs

This guide walks you through the various installation processes required to pair ROCm™ with the latest high-end AMD Radeon™ 7000 series desktop GPUs, and get started on a fully-functional environment for AI and ML development.

---

**Note:**

- MIgraphX and mGPU configuration are not currently supported by WSL
- 

### 2.2.1 Install Radeon software for WSL with ROCm

The ROCm™ Software Stack and other Radeon™ software for Windows Subsystem for Linux (WSL) components are installed using the `amdgpu-install` script to assist you in the installation of a coherent set of stack components.

- Simplifies the installation of the AMDGPU stack by encapsulating the distribution specific package installation logic and by using command line options that allows you to specify the:
  - Use case of the AMDGPU stack to be installed (WSL)
  - Combination of components (Pro stack, or user selection)
- Performs post-install checks to verify whether the installation was performed successfully.
- Installs the uninstallation script to allow you to remove the whole AMDGPU stack from the system by using a single command.

The script is provided by the installer package. See [Compatibility matrices](#) for support information.

## Prerequisites

- **WSL is installed** Ensure that WSL is installed before proceeding with ROCm installation.  
To install WSL, refer to [Windows Subsystem for Linux Documentation](#).
- **Compatible Ubuntu version is installed** Ensure that the correct Ubuntu version for the current ROCm WSL package is installed. See [Compatibility matrices](#) for support information.  
**Important!** Ensure to install ub24.04 or ub22.04 WSL container prior to installing the ROCm WSL package. For more information, refer to [Windows Subsystem for Linux Documentation](#).
- **Compatible Radeon™ Software for Windows driver is installed** WSL requires installation of the following Windows driver.  
To install the compatible driver, refer to [AMD Software: Adrenalin Edition™ 25.8.1 for WSL2](#).

## Install AMD unified driver package repositories and installer script

Select the applicable Ubuntu® version to download and install the `amdgpu-install` script on the system.

### Ubuntu 24.04

Enter the following commands to install the installer script for Ubuntu version 24.04.

```
sudo apt update
wget https://repo.radeon.com/amdgpu-install/6.4.2.1/ubuntu/noble/amdgpu-install_6.4.
˓→60402-1_all.deb
sudo apt install ./amdgpu-install_6.4.60402-1_all.deb
```

### Ubuntu 22.04

Enter the following commands to install the installer script for Ubuntu version 22.04.

```
sudo apt update
wget https://repo.radeon.com/amdgpu-install/6.4.2.1/ubuntu/jammy/amdgpu-install_6.4.
˓→60402-1_all.deb
sudo apt install ./amdgpu-install_6.4.60402-1_all.deb
```

## Install AMD unified kernel-mode GPU driver, ROCm, and graphics

After the Unified Driver Deb Package repositories are installed, run the installer script with appropriate `--usecase` parameters to install the driver components.

AMD recommends installing the WSL usecase by default.

Enter the following command to display a list of available usecases:

```
sudo amdgpu-install --list-usecase
```

### WSL usecase

AMD recommends installing the WSL usecase by default.

Run the following command to install open source graphics and ROCm.

```
amdgpu-install -y --usecase=wsl,rocm --no-dkms
```

Watch for output warning or errors indicating an unsuccessful driver installation.

**Note:** The `-y` option installs non-interactively. This step may take several minutes, depending on internet connection and system speed.

Next, *run a post-install verification check*.

### Post-install verification check

Run a post-installation check to verify that the installation is complete:

Check if the GPU is listed as an agent.

```
rocminfo
```

#### Expected result:

```
[...]
*****
Agent 2
*****
Name : gfx1100
Marketing Name: Radeon RX 7900 XTX
Vendor Name: AMD
[...]
[...]
```

### Uninstall ROCm

Run the following command to uninstall the ROCm software stack and other Radeon software for Linux components:

```
sudo amdgpu-uninstall
```

### Upgrade to newer versions of Radeon software for Linux

The recommended method to upgrade is to uninstall, followed by an install.

Radeon Software for Linux does not support in-place upgrades.

## 2.2.2 Install PyTorch for ROCm

Refer to this section for the recommended PyTorch via PIP installation method, as well as Docker-based installation.

### PCIe atomics

ROCm is an extension of HSA platform architecture, and shares queuing model, memory model, signaling and synchronization protocols.

Platform atomics are integral to perform queuing and signaling memory operations, where there may be multiple-writers across CPU and GPU agents.

For more details, see [How ROCm uses PCIe atomics](#).

### Install methods

AMD recommends the PIP install method to create a PyTorch environment when working with ROCm™ for machine learning development.

Using Docker provides portability, and access to a prebuilt Docker container that has been rigorously tested within AMD. Docker also cuts down compilation time, and should perform as expected without installation issues.

#### Option A: PyTorch via PIP installation

AMD recommends the PIP install method to create a PyTorch environment when working with ROCm™ for machine learning development.

Check [Pytorch.org](#) for latest PIP install instructions and availability. See [Compatibility matrices](#) for support information.

**Note** The latest version of Python module numpy v2.0 is incompatible with the torch wheels for this version. Downgrade to an older version is required. Example: pip3 install numpy==1.26.4

**Note** To install the following wheels, Python 3.12 must be set up.

### Install PyTorch via PIP

- Enter the following command to unpack and begin set up.

```
sudo apt install python3-pip -y
```

- Enter this command to update the pip wheel.

```
pip3 install --upgrade pip wheel
```

- Select the applicable Ubuntu version and enter the commands to install Torch and Torchvision for ROCm AMD GPU support.

This may take several minutes.

**Important!** AMD recommends proceeding with ROCm WHLs available at [repo.radeon.com](#). The ROCm WHLs available at PyTorch.org are not tested extensively by AMD as the WHLs change regularly when the nightly builds are updated.

**Important!** When manually downloading WHLs from [repo.radeon](#), ensure to select the compatible WHLs for specific Python versions.

See [Compatibility matrices](#) for support information.

### Ubuntu 24.04

```
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torch-2.6.0%2Brocm6.4.  
↪2.git76481f7c-cp312-cp312-linux_x86_64.whl  
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torchvision-0.21.0  
↪%2Brocm6.4.2.git4040d51f-cp312-cp312-linux_x86_64.whl  
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/pytorch_triton_rocm-3.  
↪2.0%2Brocm6.4.2.git7e948ebf-cp312-cp312-linux_x86_64.whl  
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torchaudio-2.6.0  
↪%2Brocm6.4.2.gitd8831425-cp312-cp312-linux_x86_64.whl  
pip3 uninstall torch torchvision pytorch-triton-rocm  
pip3 install torch-2.6.0+rocm6.4.2.git76481f7c-cp312-cp312-linux_x86_64.whl  
↪torchvision-0.21.0+rocm6.4.2.git4040d51f-cp312-cp312-linux_x86_64.whl  
↪torchaudio-2.6.0+rocm6.4.2.gitd8831425-cp312-cp312-linux_x86_64.whl pytorch_  
↪triton_rocm-3.2.0+rocm6.4.2.git7e948ebf-cp312-cp312-linux_x86_64.whl
```

#### Note

The `--break-system-packages` flag must be added when installing wheels for Python 3.12 in a non-virtual environment.

### Ubuntu 22.04

```
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torch-2.6.0%2Brocm6.4.  
↪2.git76481f7c-cp310-cp310-linux_x86_64.whl  
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torchvision-0.21.0  
↪%2Brocm6.4.2.git4040d51f-cp310-cp310-linux_x86_64.whl  
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/pytorch_triton_rocm-3.  
↪2.0%2Brocm6.4.2.git7e948ebf-cp310-cp310-linux_x86_64.whl  
wget https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/torchaudio-2.6.0  
↪%2Brocm6.4.2.gitd8831425-cp310-cp310-linux_x86_64.whl  
pip3 uninstall torch torchvision pytorch-triton-rocm  
pip3 install torch-2.6.0+rocm6.4.2.git76481f7c-cp310-cp310-linux_x86_64.whl  
↪torchvision-0.21.0+rocm6.4.2.git4040d51f-cp310-cp310-linux_x86_64.whl  
↪torchaudio-2.6.0+rocm6.4.2.gitd8831425-cp310-cp310-linux_x86_64.whl pytorch_  
↪triton_rocm-3.2.0+rocm6.4.2.git7e948ebf-cp310-cp310-linux_x86_64.whl
```

4. Update to WSL compatible runtime lib.

```
location=$(pip show torch | grep Location | awk -F ":" "\{print $2\}"  
cd ${location}/torch/lib/  
rm libhsa-runtime64.so*
```

5. **Optional step:** Using a Conda environment.

**Note** This is an optional step for users who wish to proceed with a Conda environment. AMD does not officially support and validate Conda usecases.

The libhsa-runtime64.so requires installation of GCC 12.1 at minimum. When using a Conda environment, `ImportError: version 'GLIBCXX_3.4.30' not found` is likely to occur. Upgrade GCC for Conda using the following command.

```
conda install -c conda-forge gcc=12.1.0
```

Next, *verify your PyTorch installation*.

## Option B: Docker installation

Using Docker provides portability, and access to a prebuilt Docker container that has been rigorously tested within AMD. Docker also cuts down compilation time, and should perform as expected without installation issues.

### Prerequisites to install PyTorch using Docker

Docker for Ubuntu® must be installed.

To install Docker for Ubuntu, enter the following command:

```
sudo apt install docker.io
```

### Use Docker image with pre-installed PyTorch

Follow these steps to install using a Docker image.

**Note** The latest version of Python module numpy v2.0 is incompatible with the torch wheels for this version.

Downgrade to an older version is required. Example: pip3 install numpy==1.26.4

1. Select the applicable Ubuntu version and enter the following command to pull the public PyTorch Docker image.

**Optional:** You can also download a specific and supported configuration with different user-space ROCm versions, PyTorch versions, and supported operating systems.

Refer to [hub.docker.com/r/rocm/pytorch](https://hub.docker.com/r/rocm/pytorch) to download the PyTorch Docker image.

#### Ubuntu 24.04

```
sudo docker pull rocm/pytorch:rocm6.4.2_ubuntu24.04_py3.12_pytorch_release_2.6.0
```

#### Ubuntu 22.04

```
sudo docker pull rocm/pytorch:rocm6.4.2_ubuntu22.04_py3.10_pytorch_release_2.6.0
```

2. Select the applicable Ubuntu version and start a Docker container using the downloaded image.

#### Ubuntu 24.04

```
sudo docker run -it \
--cap-add=SYS_PTRACE \
--security-opt seccomp=unconfined \
--ipc=host \
--shm-size 8G \
--device=/dev/dxg -v /usr/lib/wsl/lib/libdxcore.so:/usr/lib/libdxcore.so -v /opt/
˓ rocm/lib/libhsa-runtime64.so.1:/opt/rocm/lib/libhsa-runtime64.so.1 \
rocm/pytorch:rocm6.4.2_ubuntu24.04_py3.12_pytorch_release_2.6.0
```

### Ubuntu 22.04

```
sudo docker run -it \
--cap-add=SYS_PTRACE \
--security-opt seccomp=unconfined \
--ipc=host \
--shm-size 8G \
--device=/dev/dxg -v /usr/lib/wsl/lib/libdxcore.so:/usr/lib/libdxcore.so -v /opt/
˓→rocm/lib/libhsa-runtime64.so.1:/opt/rocm/lib/libhsa-runtime64.so.1 \
rocm/pytorch:rocm6.4.2_ubuntu22.04_py3.10_pytorch_release_2.6.0
```

This will automatically download the image if it does not exist on the host. You can also pass the `-v` argument to mount any data directories from the host onto the container.

Next, *verify the PyTorch installation*.

See PyTorch Installation for ROCm for more information.

### Verify PyTorch installation

Confirm if PyTorch is correctly installed.

1. Verify if Pytorch is installed and detecting the GPU compute device.

```
python3 -c 'import torch' 2> /dev/null && echo 'Success' || echo 'Failure'
```

**Expected result:**

```
Success
```

2. Enter command to test if the GPU is available.

```
python3 -c 'import torch; print(torch.cuda.is_available())'
```

**Expected result:**

```
True
```

3. Enter command to display installed GPU device name.

```
python3 -c "import torch; print(f'device name [0]:', torch.cuda.get_device_
˓→name(0))"
```

**Expected result:** Example: *device name [0]: Radeon RX 7900 XTX*

```
device name [0]: <Supported AMD GPU>
```

4. Enter command to display component information within the current PyTorch environment.

```
python3 -m torch.utils.collect_env
```

**Expected result:**

```
PyTorch version
ROCM used to build PyTorch
```

(continues on next page)

(continued from previous page)

OS
Is CUDA available
GPU model and configuration
HIP runtime version
MIOpen runtime version

Environment set-up is complete, and the system is ready for use with PyTorch to work with machine learning models, and algorithms.

### 2.2.3 Install ONNX Runtime for Radeon GPUs on WSL

To install ONNX Runtime on WSL, refer to [Install ONNX Runtime for Radeon GPUs](#).

**Note**

Installation instructions for ONNX Runtime on WSL are the same as Linux.

### 2.2.4 Install TensorFlow for ROCm on WSL

To install TensorFlow on WSL, refer to [Install TensorFlow for Radeon GPUs](#).

**Note**

Installation instructions for TensorFlow on WSL are the same as Linux.

### 2.2.5 Install Triton for ROCm on WSL

To install Triton on WSL, refer to [Install Triton for Radeon GPUs](#).

**Note**

Installation instructions for Triton on WSL are the same as Linux.

### 2.2.6 Install JAX for ROCm

This directory provides setup instructions and necessary files to build, test, and run JAX with ROCm support using a PIP or Docker install, suitable for both runtime and CI workflows.

**Note** These instructions are for JAX installation on **Radeon** GPUs. To install ROCm on **Instinct** GPUs, refer to [ROCM Instinct documentation](#).

### Install JAX

Follow these instructions to install JAX via PIP or Docker install.

#### PIP installation

Follow these instructions to install JAX via PIP.

**Important** The packages must be installed in the following order:

1. Install pjrt wheel.
2. Install plugin wheel.
3. Install jaxlib wheel.
4. Install jax wheel.

#### Ubuntu 24.04

Install JAX for Ubuntu 24.04.

1. Uninstall previous version

```
pip3 uninstall -y jax-rocm60-pjrt jax-rocm60-plugin jaxlib jax
```

2. Install the pjrt wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax_rocm60_pjrt-  
→0.4.35-py3-none-manylinux_2_28_x86_64.whl
```

3. Install the plugin wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax_rocm60_  
→plugin-0.4.35-cp312-cp312-manylinux_2_28_x86_64.whl
```

4. Install the jaxlib wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jaxlib-0.4.35-  
→cp312-cp312-manylinux_2_28_x86_64.whl
```

5. Install the jax wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax-0.4.35-py3-  
→none-any.whl
```

#### Ubuntu 22.04

Install JAX for Ubuntu 22.04.

1. Uninstall previous version

```
pip3 uninstall -y jax-rocm60-pjrt jax-rocm60-plugin jaxlib jax
```

2. Install the pjrt wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax_rocm60_pjrt-
→ 0.4.35-py3-none-manylinux_2_28_x86_64.whl
```

3. Install the plugin wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax_rocm60_
→ plugin-0.4.35-cp310-cp310-manylinux_2_28_x86_64.whl
```

4. Install the jaxlib wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jaxlib-0.4.35-
→ cp310-cp310-manylinux_2_28_x86_64.whl
```

5. Install the jax wheel.

```
pip install https://repo.radeon.com/rocm/manylinux/rocm-rel-6.4.2/jax-0.4.35-py3-
→ none-any.whl
```

## Docker installation

The ROCm JAX team provides prebuilt Docker images, which is the simplest way to use JAX on ROCm. These images are available on Docker Hub and come with JAX configured for ROCm.

### Ubuntu 24.04

Install JAX for Ubuntu 24.04.

1. To pull the latest ROCm JAX Docker image, run:

```
docker pull rocm/jax:rocm6.4.2-jax0.4.35-py3.12
```

**Note** For specific versions of JAX, review the periodically pushed Docker images at ROCm JAX Community on Docker Hub.

Additional Docker images are available at ROCm JAX on Docker Hub. These contain the latest ROCm version but might use an older version of JAX.

2. Once the image is downloaded, launch a container using the following command:

```
sudo docker run -it \
--cap-add=SYS_PTRACE \
--security-opt seccomp=unconfined \
--ipc=host \
--shm-size 8G \
--device=/dev/dxg -v /usr/lib/wsl/lib/libdxcore.so:/usr/lib/libdxcore.so -v /opt/
→ rocm/lib/libhsa-runtime64.so.1:/opt/rocm/lib/libhsa-runtime64.so.1 \
rocm/jax:rocm6.4.2-jax0.4.35-py3.12
```

**Note** The `--shm-size` parameter allocates shared memory for the container. Adjust it based on your system's resources if needed. Replace `$(pwd)` with the absolute path to the directory you want to mount inside the container. If you prefer to use `rocm/jax`, remember to replace `rocm/jax-community` with `rocm/jax`.

### Ubuntu 22.04

Install JAX for Ubuntu 22.04.

1. To pull the latest ROCm JAX Docker image, run:

```
docker pull rocm/jax:rocm6.4.2-jax0.4.35-py3.10
```

**Note** For specific versions of JAX, review the periodically pushed Docker images at ROCm JAX Community on Docker Hub.

Additional Docker images are available at ROCm JAX on Docker Hub. These contain the latest ROCm version but might use an older version of JAX.

2. Once the image is downloaded, launch a container using the following command:

```
sudo docker run -it \
--cap-add=SYS_PTRACE \
--security-opt seccomp=unconfined \
--ipc=host \
--shm-size 8G \
--device=/dev/dxg -v /usr/lib/wsl/lib/libdxcore.so:/usr/lib/libdxcore.so -v /opt/
˓→rocm/lib/libhsa-runtime64.so.1:/opt/rocm/lib/libhsa-runtime64.so.1 \
rocm/jax:rocm6.4.2-jax0.4.35-py3.10
```

**Note** The --shm-size parameter allocates shared memory for the container. Adjust it based on your system's resources if needed. Replace \$(pwd) with the absolute path to the directory you want to mount inside the container. If you prefer to use rocm/jax, remember to replace rocm/jax-community with rocm/jax.

### Verify installation

Refer to [Testing your JAX installation with ROCm](#) for verification steps.

## 2.2.7 Install MIGraphX for Radeon GPUs on WSL

To install MIgraphX on WSL, refer to [Install MIgraphX for Radeon GPUs](#).

#### Note

Installation instructions for MIgraphX on WSL are the same as Linux.

## USECASES

Refer to the applicable guides to optimize specific usecase performance.

### 3.1 vLLM

Refer to the applicable guides to optimize vLLM usecase performance.

#### 3.1.1 vLLM Docker image for Llama2 and Llama3

Virtual Large Language Model (vLLM) is a fast and easy-to-use library for LLM inference and serving.

Llama2 and Llama3 support is enabled via a vLLM Docker image that must be built separately (in addition to ROCm) for the current release.

For additional information, visit the [AMD vLLM GitHub page](#).

Note that this is a benchmarking demo/example. Installation for other vLLM models/configurations may differ.

#### Prerequisites

- GitHub is authenticated.

#### Additional information

- AMD recommends 40GB GPU for 70B usecases. Ensure that your GPU has enough VRAM for the chosen model.
- This example highlights use of the AMD vLLM Docker using [Llama-3 70B with GPTQ quantization](#) (as shown at Computex). However, performance is not limited to this specific Hugging Face model, and other vLLM supported models can also be used.

### Download and install Docker image

#### Download Docker image

Select the applicable Ubuntu version to download the compatible Docker image before starting.

#### Ubuntu 24.04/Python 3.12/PyTorch 2.7

```
docker pull rocm/vllm-dev:rocm6.4.2_navi_ubuntu24.04_py3.12_pytorch_2.7_vllm_0.9.2
```

#### Ubuntu 22.04/Python 3.10/PyTorch 2.7

```
docker pull rocm/vllm-dev:rocm6.4.2_navi_ubuntu22.04_py3.10_pytorch_2.7_vllm_0.9.2
```

#### Note

For more information, see [rocm/vllm-dev](#).

### Installation

Follow these steps to build a vLLM Docker image and start using Llama2 and Llama3.

1. Start the Docker container.

**Important!** Refer to *WSL specific configurations* for instructions when working in a WSL environment.

```
docker run -it \
--privileged \
--device=/dev/kfd \
--device=/dev/dri \
--network=host \
--group-add sudo \
-w /app/vllm/ \
--name <container_name> \
<image_name> \
/bin/bash
```

#### Note

- The `<container_name>` is user defined. Ensure to name your Docker using this value.

### WSL-specific configurations

**Optional: Only applicable when using a WSL configuration** Select the applicable vLLM instructions, based on your specific WSL configuration.

**Ubuntu 24.04/Python 3.12/PyTorch 2.7**

```
docker run -it \
--network=host \
--group-add=video \
--ipc=host \
--cap-add=SYS_PTRACE \
--security-opt seccomp=unconfined \
--device /dev/dxg \
--entrypoint /bin/bash \
-v /usr/lib/wsl/lib/libdxcore.so:/usr/lib/libdxcore.so \
-v /opt/rocm/lib/libhsa-runtime64.so.1:/opt/rocm/lib/libhsa-runtime64.so.1 \
-w /app/vllm/ \
--name vllm_rocm_container \
rocm/vllm-dev:rocm6.4.2_navi_ubuntu24.04_py3.12_pytorch_2.7_vllm_0.9.2 \
-c "sed -i 's/is_rocm = False/is_rocm = True/g' /opt/conda/envs/py_3.12/lib/
˓→python3.12/site-packages/vllm/platforms/__init__.py && /bin/bash"
```

**Ubuntu 22.04/Python 3.10/PyTorch 2.7**

```
docker run -it \
--network=host \
--group-add=video \
--ipc=host \
--cap-add=SYS_PTRACE \
--security-opt seccomp=unconfined \
--device /dev/dxg \
--entrypoint /bin/bash \
-v /usr/lib/wsl/lib/libdxcore.so:/usr/lib/libdxcore.so \
-v /opt/rocm/lib/libhsa-runtime64.so.1:/opt/rocm/lib/libhsa-runtime64.so.1 \
-w /app/vllm/ \
--name vllm_rocm_container \
rocm/vllm-dev:rocm6.4.2_navi_ubuntu22.04_py3.10_pytorch_2.7_vllm_0.9.2 \
-c "sed -i 's/is_rocm = False/is_rocm = True/g' /opt/conda/envs/py_3.10/lib/
˓→python3.10/site-packages/vllm/platforms/__init__.py && /bin/bash"
```

- Clone the Hugging Face GitHub repository within the Docker container.

```
apt update
apt install git-lfs
git lfs clone https://huggingface.co/TechxGenus/Meta-Llama-3-70B-Instruct-GPTQ
```

- Run benchmarks with the Docker container.

```
python3 /app/vllm/benchmarks/benchmark_latency.py --model /app/vllm/Meta-Llama-3-
˓→70B-Instruct-GPTQ -q gptq --batch-size 1 --input-len 1024 --output-len 1024 --
˓→max-model-len 2048
```

**Note** Ensure that the model is downloaded and vLLM checkout is set to your current directory within the container described in Step 3.

**Note** Set the following environment variable prior to running models using vLLM in V1 mode:  
**VLLM\_V1\_USE\_PREFILL\_DECODE\_ATTENTION=1**

Your environment is set up to use Llama2 and Llama3.

### 3.1.2 GEMM tuning for model inferencing with vLLM

**Note** Tuning must be done for specific tensor-parallel-size.

#### Collect GEMM shape details

Collect GEMM shape details used in this model inference, and ensure `tp` is specified based on targeted setup accordingly (split over N GPUs).

```
VLLM_TUNE_GEMM=1 VLLM_UNTUNE_FILE=untuned_gemm.csv python <vllm_path>/benchmarks/
˓→benchmark_throughput.py --model <model_path> --trust-remote-code --dataset <dataset_
˓→path>/ShareGPT_V3_unfiltered_cleaned_split.json --num_prompts 1000 [--distributed_
˓→executor_backend mp --tensor-parallel-size N]
```

#### Conduct GEMM tuning

The `indtype/outdtype` must be correctly specified, and aligned with `dtype` used for model inferencing.

For quantized int4 models, use f16 or bf16. Refer to [ROCm vLLM Github](#) (Sep 23, 2024, included in v0.6.1.post1+rocm release) as a starting point.

The generated `untuned_gemm.csv` contains `dtype` info that is used by default, and is only overridden by `indtype` and `outdtype` when specified.

**Important** Do not use the `tp` option, as it is designed for other usecases.

**Note** If `gradlib` fails with OOM then set `CACHE_INVALIDATE_BUFFERS` to a primary number (such as 11 or 7 or 3 or even 1).

```
python <vllm_path>/gradlib/gradlib/gemm_tuner.py --input_file untuned_gemm.csv --
˓→tuned_file tuned_gemm_tpN.csv [--indtype f16 --outdtype f16]
```

#### Run vLLM inference with tuned GEMM

Enter the following command.

```
VLLM_TUNE_FILE=tuned_gemm_tpN.csv python <vllm_path>/benchmarks/benchmark_throughput .
˓→py --model <model_path> --trust-remote-code --dataset <dataset_path>/ShareGPT_V3_
˓→unfiltered_cleaned_split.json --num_prompts 1000 [--distributed_executor_backend mp_
˓→--tensor-parallel-size N]
```

## 3.2 ComfyUI

Install ComfyUI and its extensions on your system.

### 3.2.1 Install ComfyUI and MIGraphX extension

Install ComfyUI and the MIGraphX extension on AMD GPUs.

#### Prerequisites

- ROCm is installed.

#### Installation

Follow these steps to install ComfyUI.

1. Install the Python venv package for the applicable Python version.

##### Python 3.12/Ubuntu 24.04

```
sudo apt install python3.12-venv
```

##### Python 3.10/Ubuntu 22.04

```
sudo apt install python3.10-venv
```

2. Create a Python virtual environment.

```
python3 -m venv venv --system-site-packages  
source venv/bin/activate
```

3. Install PyTorch (ROCM) wheels in a virtual environment.

Install the latest PyTorch ROCm wheels in the environment created.

#### Note

See [Install PyTorch for Radeon GPUs](#) for install instructions.

4. Clone the ComfyUI repository.

```
git clone https://github.com/comfyanonymous/ComfyUI.git
```

#### Note

See [ComfyUI Github](#) for more information.

5. Modify requirements.txt.

To avoid conflicts with ROCm's PyTorch packages, comment out the following packages in requirements.txt.

- torch
- torchaudio
- torchvision

```
cd ComfyUI  
sudo nano requirements.txt
```

### Post-edit example

```
comfyui-frontend-package==1.19.9
comfyui-workflow-templates==0.1.14
#torch
torchsde
#torchvision
#torchaudio
numpy>=1.25.0
einops
transformers>=4.28.1
tokenizers>=0.13.3
sentencepiece
safetensors>=0.4.2
aiohttp>=3.11.8
yarl>=1.18.0
pyyaml
Pillow
scipy
tqdm
psutil
#non essential dependencies:
kornia>=0.7.1
spandrel
soundfile
av>=14.2.0
pydantic~2.0
```

6. Install ComfyUI requirements.

```
pip install -r requirements.txt
```

7. Load Stable Diffusion 3 model.

Load a Stable Diffusion 3 model into the following directory:

```
cd ComfyUI/models/checkpoints
```

#### Note

As an example, the `sd3_medium_incl_clips.safetensors` model can be added into this directory.

8. Launch ComfyUI.

#### Note

Users must return to the home directory of the repository in order to run `main.py`.

```
python3 main.py
```

## Installing the MIGraphX node for ComfyUI

The MIGraphX Node for ComfyUI offers accelerated inference on Stable Diffusion 3 models.

### Note

Refer to [ComfyUI MIGraphX Github](#) for more information.

1. Install the MIGraphX node.

```
cd ComfyUI/custom_nodes
git clone https://github.com/pnikolic-amd/ComfyUI_MIGraphX.git
cd ComfyUI_MIGraphX
pip install -r requirements.txt
```

2. Launch ComfyUI.

### Note

The ROCm path to Python must be exported to prevent a module missing error for MIGraphX.

```
export PYTHONPATH=/opt/rocm-X.X.X/lib:$PYTHONPATH
```

Replace the ROCm version with your applicable version. For example, /opt/rocm-6.4.1/lib.

```
cd ComfyUI
# For the best performance, run the following environment variable:
export MIGRAPHX_MLIR_USE_SPECIFIC_OPS="attention"
python3 main.py
```

## Using the MIGraphX node for ComfyUI

The MIGraphX Node for ComfyUI includes ready-to-use templates, each preconfigured with optimized settings for the supported Stable Diffusion models.

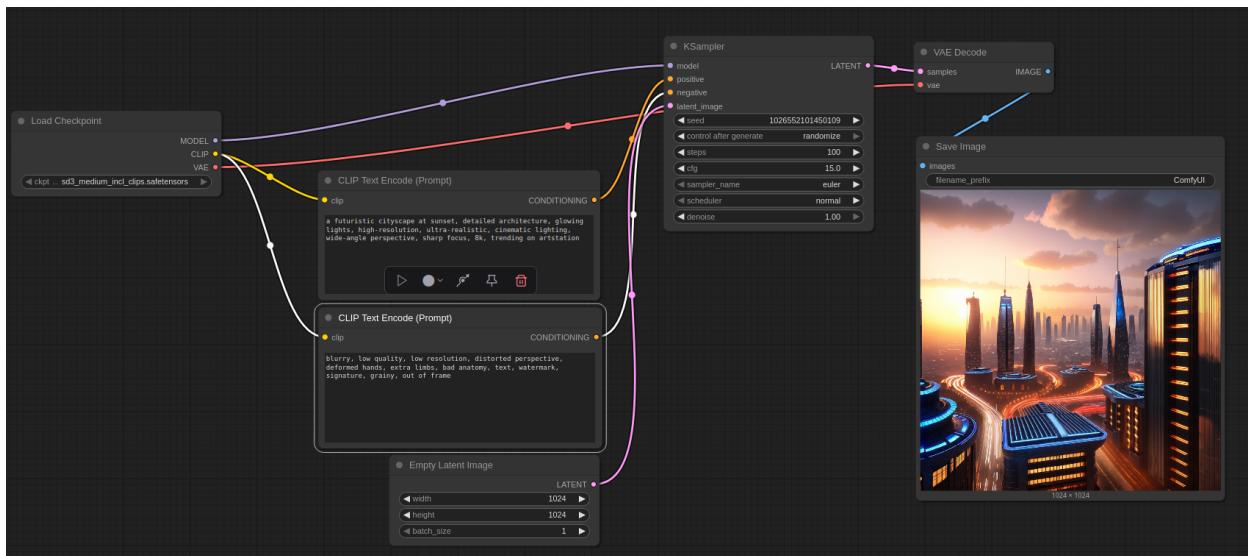
- To access them, go to: Workflow > Browse Templates > ComfyUI\_MIGraphX
- Each template is optimized for a specific Stable Diffusion model. Be sure to load the corresponding model before using the template.

## Example usage

### Note

The initial run takes some time to allow MIGraphX to warm up and tune the model.

If the recommended `sd3_medium_incl_clips.safetensors` is loaded, the `SD3M_example_PYTORCH` template can be selected.



## COMPATIBILITY MATRICES

This section provides information on the compatibility of ROCm™ components, Radeon™ GPUs, and the Radeon Software for Linux® version (Kernel Fusion Driver) and Windows Subsystem for Linux (WSL).

---

**Note:** To rollback support matrices and install instructions for previous versions, click **Version List** located at the top-right corner of the screen, or select the version (v:) menu on the bottom-left.

---

Linux

WSL

*Linux Compatibility*

*WSL Compatibility*

### 4.1 Linux support matrices by ROCm version

#### 4.1.1 ROCm 6.4.2

Compatible OS, GPU, and framework support matrices for the latest ROCm release.

To rollback support matrices and install instructions for previous versions, click **Version List** located at the top-right corner of the screen, or select the version (v:) menu on the bottom-left.

##### OS support matrix

OS	Kernel	Supported
Ubuntu® 24.04.2 Desktop Version with HWE	Ubuntu kernel 6.11	Yes
Ubuntu® 22.04.5 Desktop Version with HWE	Ubuntu kernel 6.8	Yes
RHEL 9.6	Linux kernel 5.14	Yes

**Note** AMD currently provides preliminary support for Ubuntu 24.04.3 using the Ubuntu 24.04.2 installer.

## GPU support matrix

ROC Ver-	Radeon™ Version	Supported AMD Radeon™ Hardware
Ver-	Soft-	ware for
sion	Linux®	
6.4.2	22.04	AMD Radeon RX 9070AMD Radeon RX 9070 XTAMD Radeon RX 9070 GREAMD Radeon
	24.04	AI PRO R9700AMD Radeon RX 9060AMD Radeon RX 9060 XTAMD Radeon RX 7900 XTX-
	25.10	AMD Radeon RX 7900 XTAMD Radeon RX 7900 GREAMD Radeon PRO W7900AMD Radeon PRO W7900 Dual SlotAMD Radeon PRO W7800AMD Radeon PRO W7800 48GBAMD Radeon RX 7800 XTAMD Radeon PRO W7700AMD Radeon RX 7700 XT

## Framework + ROCm support matrices

View the ROCm support matrices for PyTorch, ONNX, and TensorFlow frameworks.

### PyTorch + ROCm support matrix

PyTorch Version	ROCM Version	Comments
2.6.0	6.4.2	Official production support. See <a href="#">Install PyTorch for Radeon GPUs</a> .
2.8/Nightly	6.4	Available from PyTorch.org nightly builds, not tested extensively by AMD.
2.7/Stable	6.4	Not supported for Radeon 7000 series.

#### AI Data Types

- FP32
- FP16
- Mixed precision (FP32/FP16)
- INT8

### ONNX Runtime + ROCm support matrix

ONNX RT Version	ROCM Version	Comments
1.21	6.4.2	Official production support. See <a href="#">Install ONNX for Radeon GPUs</a> .

#### AI Data Types

- FP32
- FP16
- INT8 (MIGraphX)
- Mixed precision (FP32/FP16)

**Note** Refer to [Installation Instructions to Get Started with ONNX Runtime](#) for more information.

## TensorFlow + ROCm support matrix

TensorFlow Version	ROCM Version	Comments
2.18.1	6.4.2	Official production support. See <a href="#">Install TensorFlow for Radeon GPUs</a> .

### AI Data Types

- FP32
- FP16

## Triton + ROCm support matrix

Triton Version	ROCM Version	Comments
3.0.0	6.4.2	Official production support. See <a href="#">Install Triton for Radeon GPUs</a> .

**Note** Refer to the official [Triton documentation](#) for more information.

### 4.1.2 Docker support matrix

See Docker Image Support Matrix for the latest version of the software support matrices for ROCm container releases.

## 4.2 WSL support matrices by ROCm version

### 4.2.1 ROCm 6.4.2

This section provides information on the compatibility of ROCm™ components, Radeon™ GPUs, and the Radeon Software for Windows Subsystem for Linux® (WSL).

To rollback support matrices and install instructions for previous versions, click **Version List** located at the top-right corner of the screen, or select the version (v.) menu on the bottom-left.

### OS support matrix

OS	Kernel	Supported
Ubuntu® 24.04.2 Desktop Version with HWE	WSL2-Linux-Kernel 5.15	Yes
Ubuntu 22.04 or Ubuntu 22.04 LTS	WSL2-Linux-Kernel 5.15	Yes

**Note** Refer to [How to install Linux on Windows with WSL](#) for up-to-date OS compatibility information.

## GPU support matrix

ROC Version	Radeon™ Software Version	Radeon™ Version	Supported AMD Radeon™ Hardware
6.4.2	22.04	AMD Software: Edition™ 25.8.1 for WSL	AMD Radeon RX 9070AMD Radeon RX 9070 XTAMD Radeon RX 9070 GREAMD Radeon AI PRO R9700AMD Radeon RX 9060AMD Radeon RX 9060 XTAMD Radeon RX 7900 XTXAMD Radeon RX 7900 XTAMD Radeon RX 7900 GREAMD Radeon PRO W7900AMD Radeon PRO W7900 Dual SlotAMD Radeon PRO W7800AMD Radeon PRO W7800 48GBAMD Radeon RX 7800 XTAMD Radeon PRO W7700
	24.0425.1	Adrenalin Edition™ 25.8.1 for WSL2	

## Framework + ROCm support matrices

View the ROCm support matrices for PyTorch, ONNX, and TensorFlow frameworks.

### PyTorch + ROCm support matrix

PyTorch Version	ROCM Version	Comments
2.6.0	6.4.2	Official production support. See <a href="#">Install PyTorch for Radeon GPUs on WSL</a> .
2.8/Nightly	6.4	Available from PyTorch.org nightly builds, not tested extensively by AMD.
2.7/Stable	6.4	Not supported for Radeon 7000 series.

#### AI Data Types

- FP32
- FP16
- Mixed precision (FP32/FP16)
- INT8

### ONNX + ROCm support matrix

ONNX Version	ROCM Version	Comments
1.21	6.4.2	Official production support. See <a href="#">Install ONNX for Radeon GPUs on WSL</a> .

#### AI Data Types

- FP32
- FP16
- INT8 (MIGraphX)

- Mixed precision (FP32/FP16)

**Note** Refer to [Installation Instructions to Get Started with ONNX Runtime](#) for more information.

### TensorFlow + ROCm support matrix

TensorFlow Version	Ver-	ROCm Version	Ver-	Comments
2.18.1		6.4.2		Official production support. See <a href="#">Install TensorFlow for Radeon GPUs on WSL</a> .

#### AI Data Types

- FP32
- FP16

### Triton + ROCm support matrix

Triton Version	ROCm Version	Comments
3.0.0	6.4.2	Official production support. See <a href="#">Install Triton for Radeon GPUs on WSL</a> .

**Note** Refer to the official [Triton documentation](#) for more information.

### 4.2.2 Docker support matrix

See [Docker Image Support Matrix](#) for the latest version of the software support matrices for ROCm container releases.



## LIMITATIONS AND RECOMMENDED SETTINGS

This section provides information on software and configuration limitations.

**Note** For ROCm on Instinct known issues, refer to [AMD ROCm Documentation](#)

For OpenMPI limitations, see [ROCm UCX OpenMPI on Github](#)

### 5.1 6.4.2 release known issues

- Intermittent script failure may be observed while running Stable Diffusion workloads with ONNX Runtime and MIGraphX. Users experiencing this issue are recommended to use MXR files instead of rebuilding the model as a temporary workaround.
- Intermittent system or application crash may be observed while running Luxmark in conjunction with other compute workloads on Radeon™ RX 9000 series graphics products. Users experiencing this issue are recommended to shut down other applications and workloads while running Luxmark.
- Intermittent script failure (out of memory) may be observed while running high-memory LLM workloads with multiple GPUs on Radeon RX 9060 series graphics products.
- Intermittent script failure may be observed while running BERT training workloads with JAX.
- Intermittent script failure may be observed while running Stable Diffusion 2.1 FP16 workloads with JAX.
- Intermittent script failure (out of memory) may be observed while running Llama2 FP16 workloads with ONNX Runtime and MIGraphX.
- Intermittent system or application crash may be observed while running TensorFlow ResNet50 training workloads.
- Increased memory consumption may be observed while running TensorFlow Resnet50 training workloads.
- Lower than expected performance may be observed while running Llama 3 8B inference workloads with Llama.cpp

#### 5.1.1 WSL specific issues

- Intermittent script failure may be observed while running Llama 3 inference workloads with vLLM in WSL2. End users experiencing this issue are recommended to follow vLLM setup instructions [here](#).
- Intermittent script failure or driver timeout may be observed while running Stable Diffusion 3 inference workloads with JAX.
- Lower than expected performance may be observed while running inference workloads with JAX in WSL2.
- Intermittent script failure may be observed while running Resnet50, BERT, or InceptionV3 training workloads with ONNX runtime.

- Output error message (resource leak) may be observed while running Llama 3.2 workloads with vLLM.
- Output error message (VaMgr) may be observed while running PyTorch workloads in WSL2.
- Intermittent script failure or driver timeout may be observed while running Stable Diffusion inference workloads with TensorFlow.
- Intermittent application crash may be observed while running Stable Diffusion workloads with ComfyUI and MI-GraphX on Radeon™ RX 9060 series graphics products.
- Intermittent script failure may occur while running Stable Diffusion 2 workloads with PyTorch and MIGraphX
- Intermittent script failure may occur while running LLM workloads with PyTorch on Radeon™ PRO W7700 graphics products.
- Lower than expected performance (compared to native Linux) may be observed while running inference workloads (eg. Llama2, BERT) in WSL2.

**Important!** Radeon™ PRO Series graphics cards are not designed nor recommended for datacenter usage. Use in a datacenter setting may adversely affect manageability, efficiency, reliability, and/or performance. GD-239.

**Important!** ROCm is not officially supported on any mobile SKUs.

## 5.2 Multi-GPU configuration

AMD has identified common errors when running ROCm™ on Radeon™ multi-GPU configuration at this time, along with the applicable recommendations.

See [mGPU known issues and limitations](#) for a complete list of mGPU known issues and limitations.

## 5.3 Windows Subsystem for Linux (WSL)

WSL recommended settings and limitations.

### 5.3.1 WSL recommended settings

**Optimizing GPU utilization** WSL overhead is a noted bottleneck for GPU utilization. Increasing the batch size of operations will load the GPU more optimally, reducing time required for AI workloads. Optimal batch sizes vary by model, and macro-parameters.

### 5.3.2 ROCm support in WSL environments

#### ROCM-smi support

Due to WSL architectural limitations for native Linux User Kernel Interface (UKI), rocm-smi is not supported.

Issue	Limitations
UKI does not currently support rocm-smi	<b>No current support for:</b> Active compute processes GPU utilization Modifiable state features

### ROCM-profiler support

Not currently supported.

### Debugger

Not currently supported.

### 5.3.3 Running PyTorch in virtual environments

Running PyTorch in virtual environments requires a manual `libhsa-runtime64.so` update.

When using the WSL usecase and `hsa-runtime-rocr4wsl-amdgpu` package (installed with PyTorch wheels), users are required to update to a WSL compatible runtime lib.

#### Solution:

Enter the following commands:

```
location=$(pip show torch | grep Location | awk -F ":" '{print $2}')
cd ${location}/torch/lib/
rm libhsa-runtime64.so*
cp /opt/rocm/lib/libhsa-runtime64.so.1.2 libhsa-runtime64.so
```



---

**CHAPTER  
SIX**

---

**AI COMMUNITY**

Want to share your experiences, find answers, or contribute to resolving issues?

Explore the AMD [AI Community Forum](#), where you will find a like-minded community, passionate about all things AI!



---

**CHAPTER  
SEVEN**

---

**REPORT A BUG**

Found a defect? Report issues through [ROCM GitHub Issues](#), and contribute to improving our user experience.